

Zero Trust, Software Defined Perimeter, and P4


Open Networking Summit, Europe 2019

Omer Anson
Huawei

Introduction

Who Am I?

- Omer Anson
- Software Physicist
 - About 11 years under the keyboard
 - Almost four years at Huawei
 - Working on (mostly):
 - Linux
 - Networking
 - Cloud



Introduction

Motivation

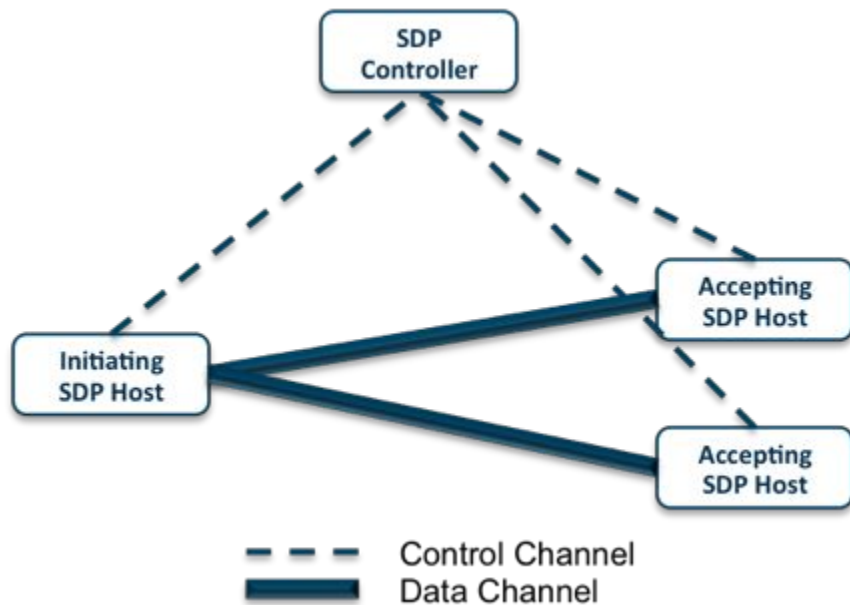
Zero-Trust Networking

- Everyone and everything has an identity
- Doesn't matter who they are
 - Lambda
 - Microservices
 - Users
 - External entities
 - PaaS
 - IaaS
- Pure Whitelist Security
 - Only permitted entities can communicate.
- Helps prevent data breaches
 - Read: Saves money!



Software Defined Perimeter (SDP)

- VPN mesh with whitelist policy
- All endpoints are authenticated and identified
- Endpoints only responded to permitted endpoints



Introduction

Technologies

netfilter

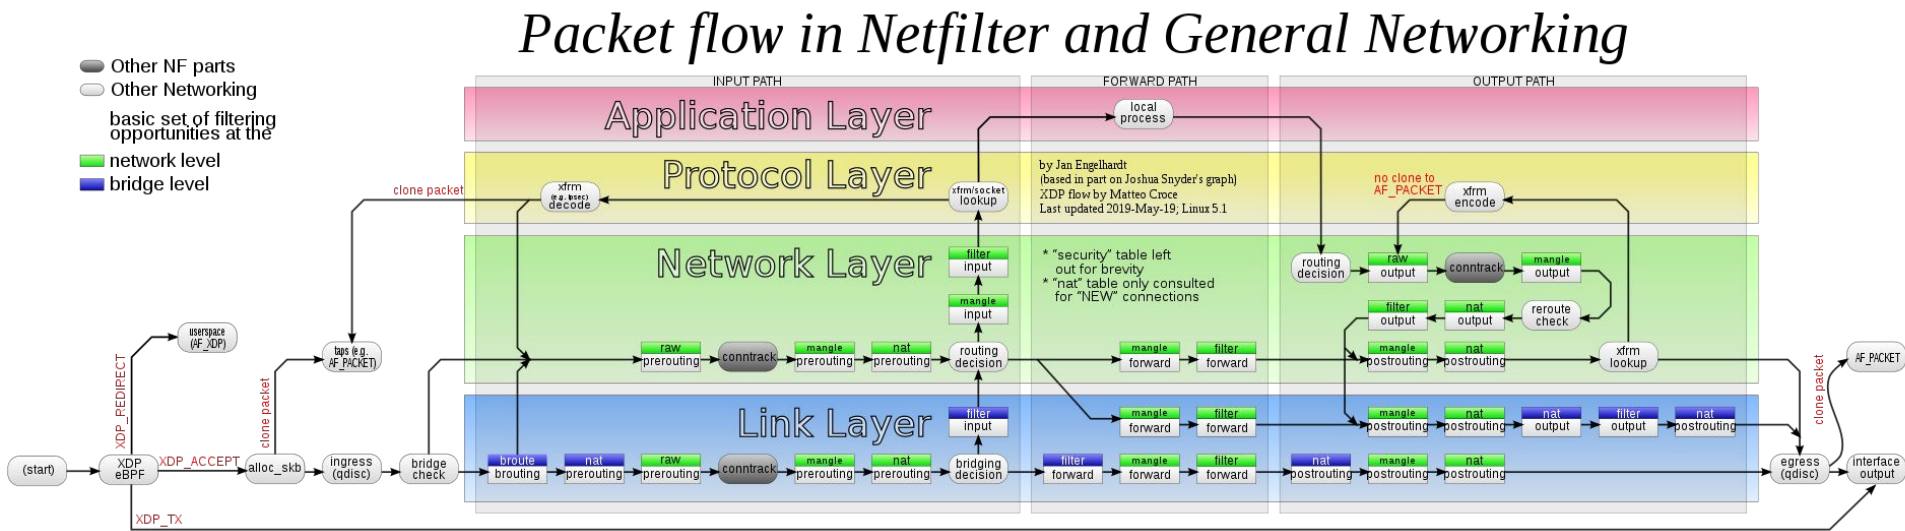



Image from Wikipedia: <https://en.wikipedia.org/wiki/Iptables>

OVS / OpenFlow



- OpenVSwitch - Virtual OpenFlow switch
- Allows to define whole pipelines
- Match-action rules (flows)
- Example flows:

```
table=55, priority=100, metadata=0x1f, dl_dst=fa:16:3e:d3:3d:8b actions=load:0x66->NXM_NX_REG7[], resubmit(,75)
```

```
table=55, priority=100, metadata=0x1f, dl_dst=fa:16:3e:c6:23:3f actions=load:0x67->NXM_NX_REG7[], resubmit(,75)
```

```
table=55, priority=200, metadata=0x1f, dl_dst=fa:16:3e:2e:a1:eb actions=load:0x10->NXM_NX_REG5[], resubmit(,60)
```

```
table=55, priority=200, metadata=0x1f, dl_dst=fa:16:3e:3a:39:bf actions=load:0x10->NXM_NX_REG5[], resubmit(,60)
```

```
table=55, priority=100, metadata=0x1f, dl_dst=01:00:00:00:00:00/01:00:00:00:00:00
```

```
>>>>>>>actions=load:0x67->NXM_NX_REG7[], resubmit(,75), load:0x66->NXM_NX_REG7[], resubmit(,75), load:0->NXM_NX_REG7[], resubmit(,75)
```

extended Berkeley Packet Filter (eBPF)

- “Super powers have finally come to Linux”

- Brendan Gregg

- In-kernel Virtual Machine
- Userspace code running in kernel
 - Change behaviour during runtime
 - In our case: Packet processing
 - Provably safe
- Many entry-points
 - Device and socket packet filtering

Why Not Write Directly in eBPF?

- Lots of boilerplate
 - Redundant verification
 - Can be inferred automatically from P4 code
 - Map structure must be defined beforehand
 - Can be inferred automatically from P4 code
 - Internal test, without defining headers:
 - P4: 70 LOC
 - eBPF: 160 LOC



P4

- From the website:
 - P4 programs specify how a switch processes packets.
- Domain specific language
- Specifies how to process packets
- Protocol Independent
- Compiled to different backends and architectures

```
// Parse packet
parser ingress_parser(packet_in packet, out headers_t hdr,
                      inout metadata_t meta, inout standard_metadata_t std_meta) {
    state start {
        packet.extract(hdr.ethernet);
        // ...
    }
}
```


P4

- Missing compilation options
- P4 has a reference implementation switch, **but**
 - Slow
 - Userspace
 - Geared towards flexibility, not speed
 - Unstable
 - At least, compared to the Linux kernel
- So let's compile to eBPF!



But there **is** a P4->eBPF compiler!

- Current eBPF compilation is feature-poor
 - Can't change packet structure
 - No redirection
 - Packets either pass or drop
- Translates to C
 - High level language
 - Optimisations are less precise
 - Another layer of indirection
 - Programme structure information loss
 - More code; more room for bugs

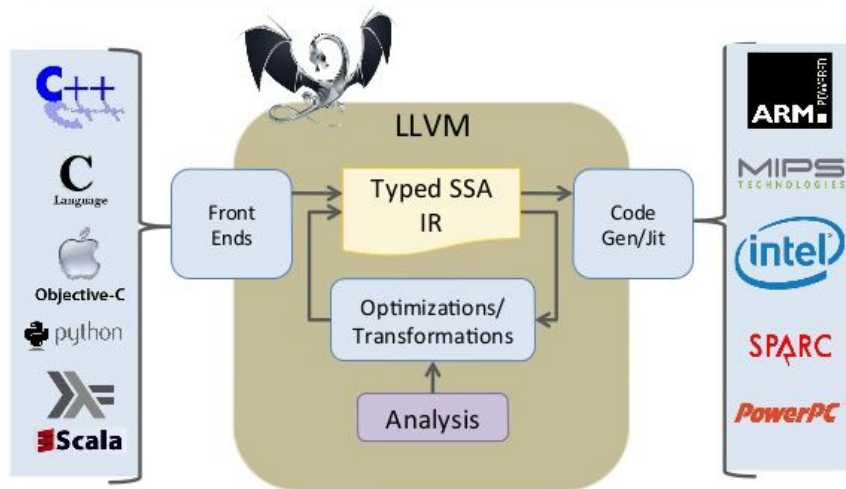


So Let's Write a P4 Compiler

LLVM

LLVM Compiler Infrastructure


[Lattner et al.]



- Compiler and toolchain technologies
- Modular and reusable

Skip the Middleman

- Translate Directly to LLVM IR
 - Better support for optimisation
 - Add entry-points and placeholders
 - Example: Strong use of *mem2reg*
 - Example: undef -> 0 replacement
 - Better control over generated code
 - Can add debug symbols and code locators
 - Less computation
 - Meaning less room for bugs



The Technicalities

The Technicalities

- So we wrote our code in P4
 - Implemented identity
 - For Zero-Trust
 - Implemented security
 - For Zero-Trust and SDP
 - Implemented the kitchen sink
- Now what?

The Technicalities

Code


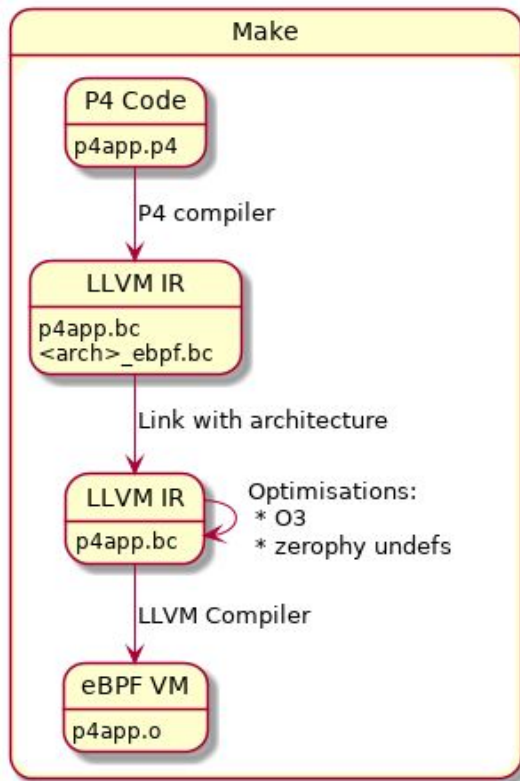
The Technicalities - Code (L2)

```
1 control l2(inout headers_t hdr, inout metadata_t meta, inout standard_metadata_t istd) {
2     action set_dst_lport(bit<32> port) {
3         meta.dst_lport = port;
4     }
5     table l2_forward_tbl {
6         key = {
7             hdr.ethernet.dstAddr: exact;
8         }
9         actions = {
10             set_dst_lport;
11         }
12     }
13
14     apply {
15         l2_forward_tbl.apply();
16     }
17 }
```

The Technicalities

Compilation

Compilation




The Technicalities

Loading (an example)


“Physical” Network Structure

- Single P4 Controller namespace
- Each network element in its own namespace
 - Connected with a veth pair
- Where to load compiler output?
 - P4 Controller only
 - Leg of veth pair
- **We're good to go!**



You Promised Zero-Trust and SDP!

- Write the logic in P4
 - For Zero-Trust
 - For SDP
 - For the pygmy marmoset passing the packets
- Compile it
 - That's where this work comes in
- Load it in the gateway



Summary

Conclusion

- Wrote a Zero-Trust and SDP Gateway in P4
- Wrote a compiler: P4 -> LLVM
- Use LLVM optimisations
- It works. Managed to compile, run, and test

Thank You!

Questions?

Thank You!

(Come Again)