

# 开源边缘计算平台

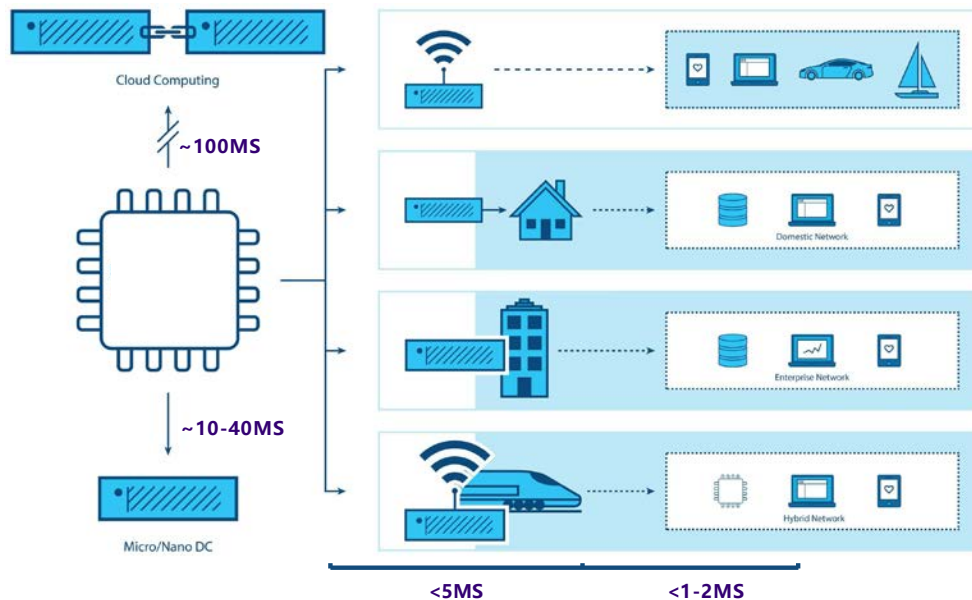
## Introduce The Open Source Edge Computing Platform - StarlingX

汇报人：九州云信息科技有限公司  
汇报日期：2018年11月


# What Is Driving Edge Computing?

- A. Latency
- B. Bandwidth
- C. Security
- D. Connectivity

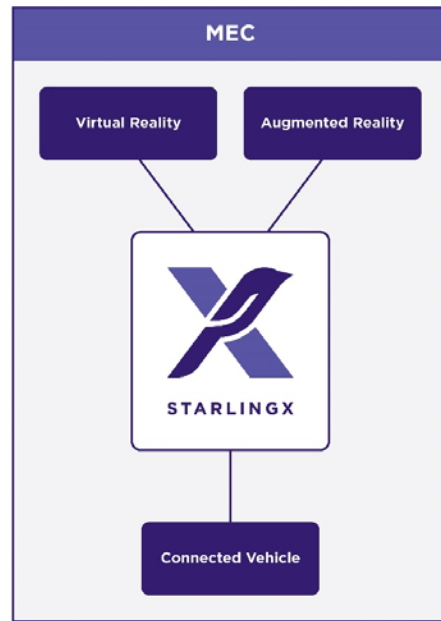
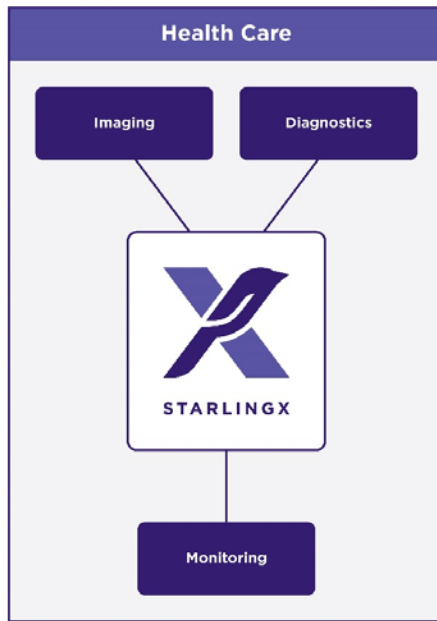
“WHERE” MATTERS



# Edge Computing Use Cases

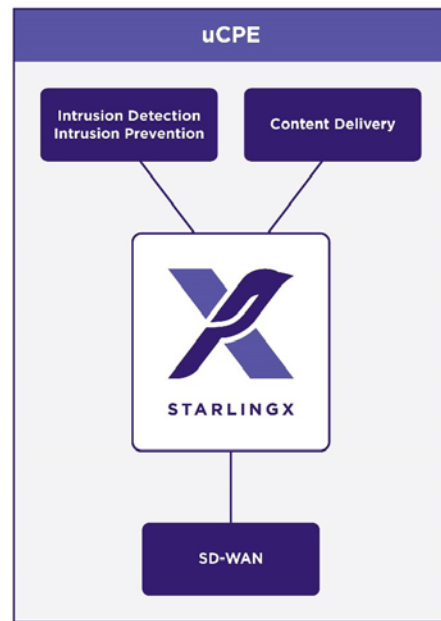
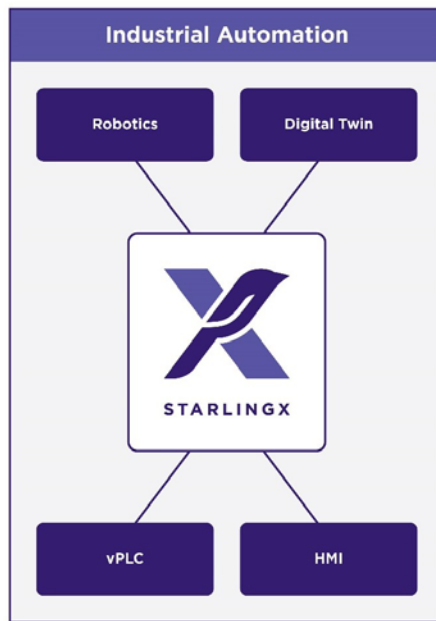
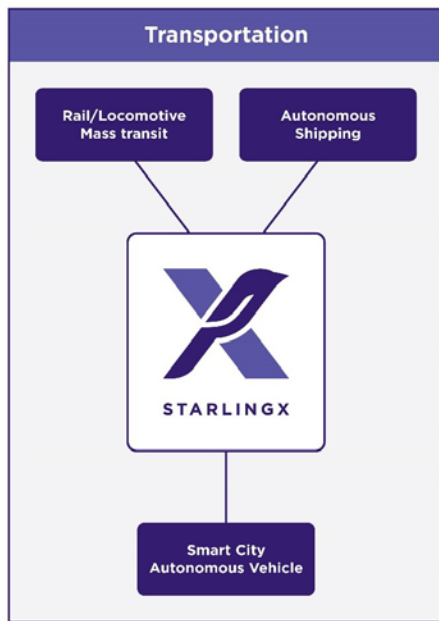


*vRAN == virtual Radio Access Network*



*MEC == Multi-access Edge Computing*

# Edge Computing Use Cases



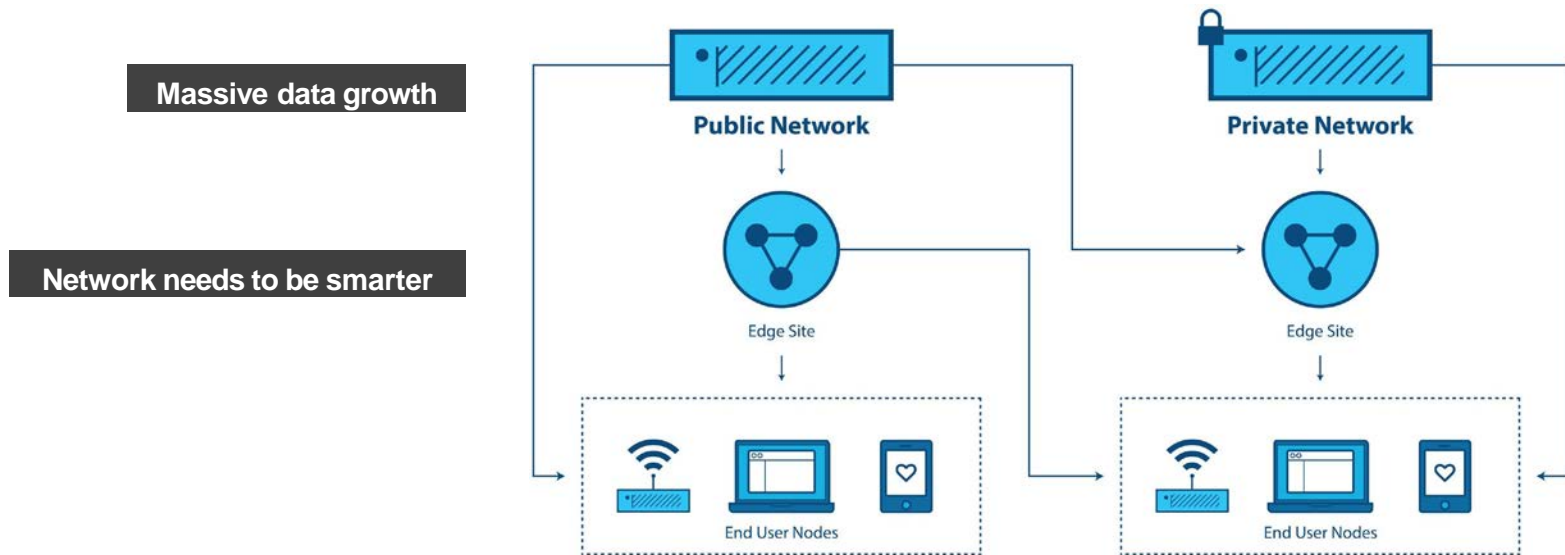
*uCPE == universal Customer Premises Equipment*

# Solving The Challenges: Edge Initiatives





# What Problems Is StarlingX Solving?



1. Distributed infrastructure demands a different architecture
2. The maturity and robustness of Cloud is required everywhere
3. Managing a massively distributed compute environment is hard

# Intent of the StarlingX Project

## Re-Configure Proven Cloud Technologies for Edge Compute

- Orchestrate system-wide
- Deploy and manage Edge clouds, share configurations
- Simplify deployment to geographically dispersed, remote Edge regions



Transportation	Energy
Manufacturing	Retail
Video	Smart cities
Healthcare	Drones

# StarlingX Technology



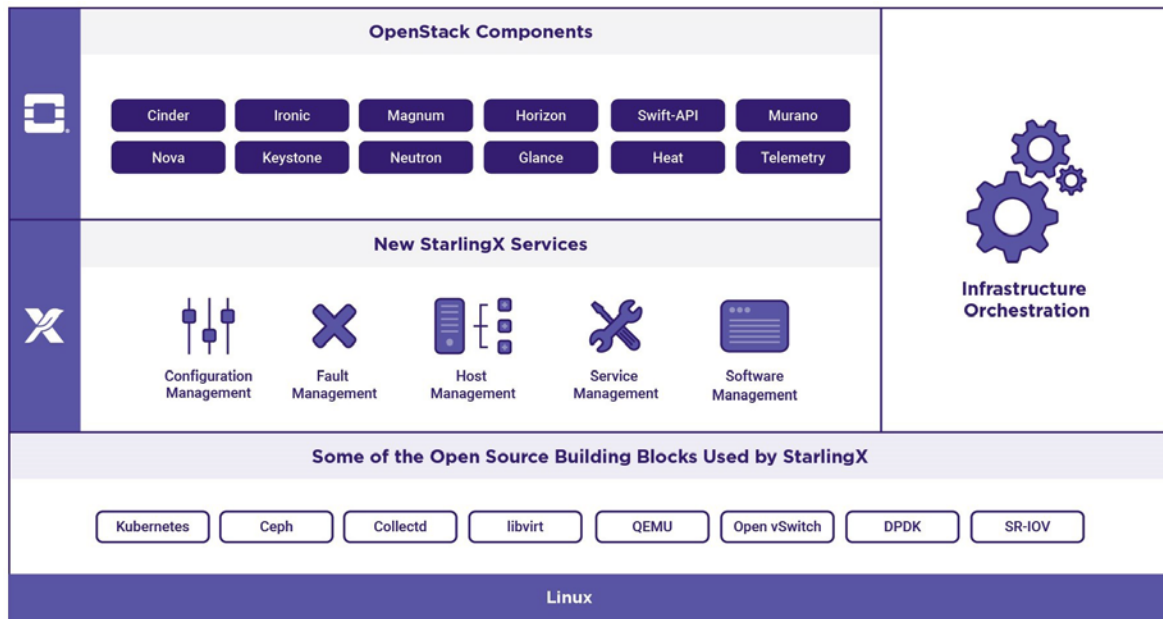
# StarlingX – Edge Virtualization Platform

StarlingX provides a **deployment-ready, scalable, highly reliable** Edge infrastructure software platform

Services from the StarlingX virtualization platform focus on

- **Easy deployment**
- **Low touch manageability**
- **Rapid response to events**
- **Fast recovery**

Think control at the Edge, control between IoT and Cloud, control over your virtual machines.



# Scalability from Small to Large

## Single Server

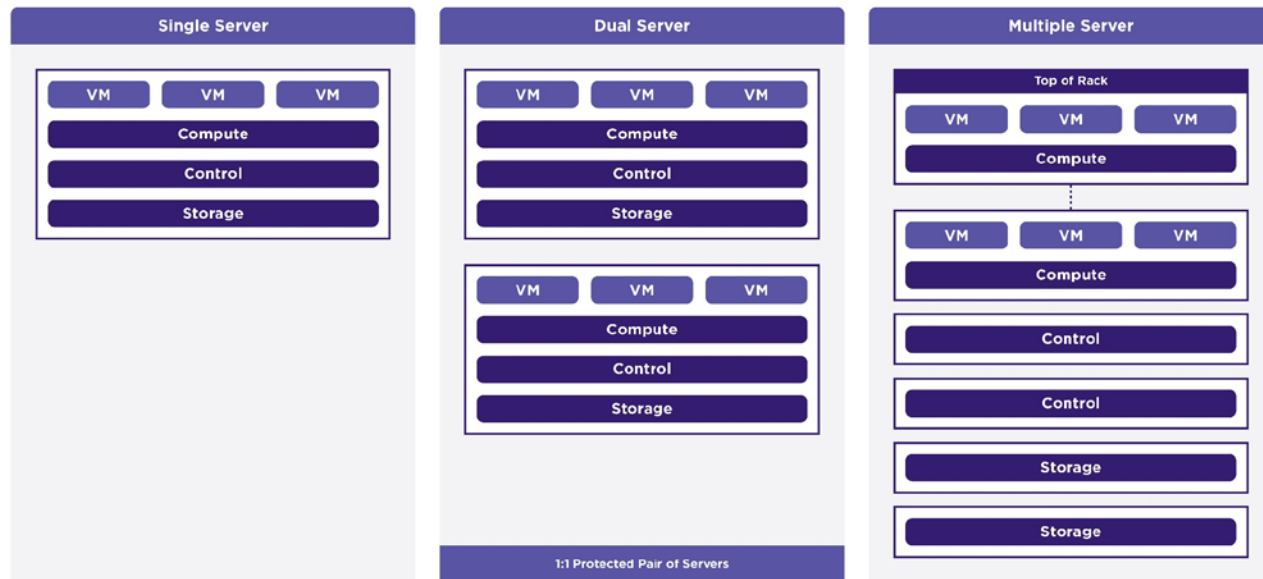
- Runs all functions

## Dual Server

- Redundant design

## Multiple Server

- Fully resilient and geographically distributable



STARLINGX

Physical Server ☐

# Configuration Management

## Manages installation

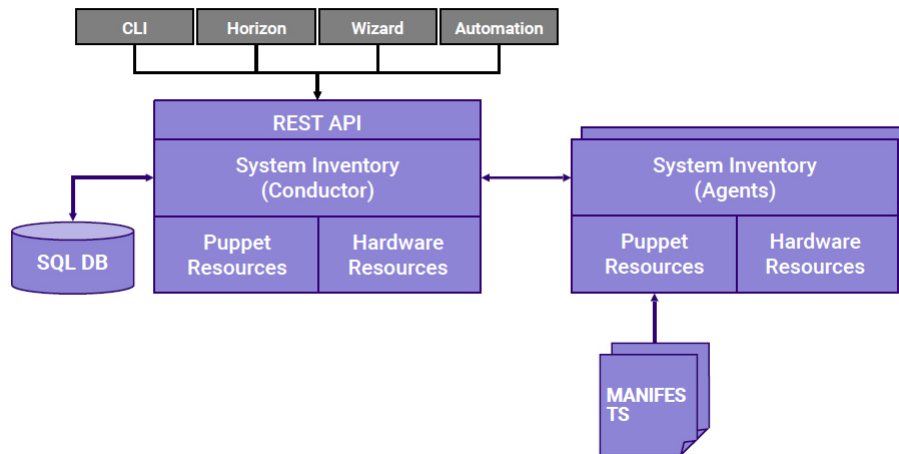
- Auto-discover new nodes
- Manage installation parameters (i.e. console, root disks)
- Bulk provisioning of nodes through XML file

## Nodal Configuration

- Node role, role profiles
- Core, memory (including huge page)
- assignments
- Network Interfaces and storage assignments

## Inventory Discovery

- CPU/cores, SMT, processors, memory, huge pages
- Storage, ports
- GPUs, storage, Crypto/compression H/W



System Configuration and Setup

# Host Management

**Full life-cycle management of the host**

**Detects and automatically handles host failures and initiates recovery**

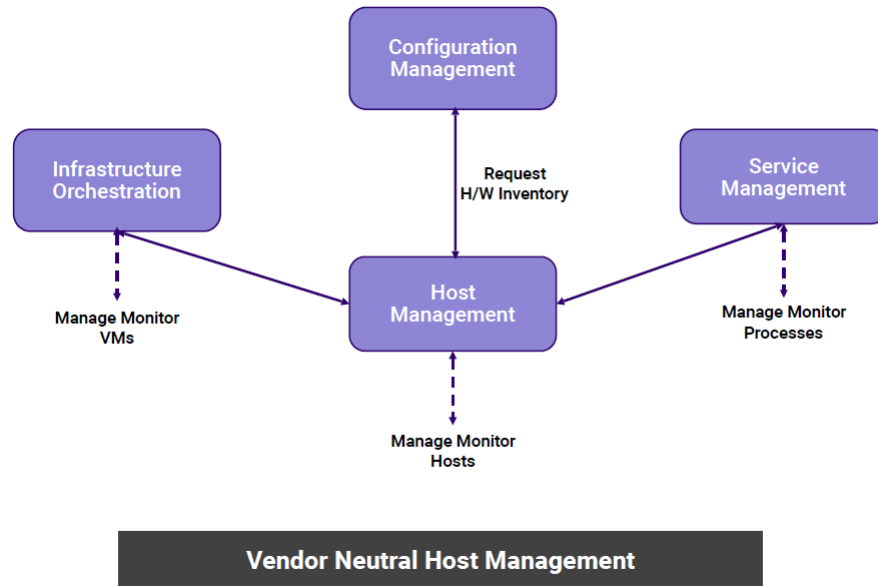
**Monitoring and alarms for**

- Cluster connectivity, critical process
- failures
- Resource utilization thresholds, interface
- states
- H/W fault / sensors, host watchdog
- Activity progress reporting

**Interfaces with board management (BMC)**

- For out of band reset
- Power-on/off
- H/W sensor monitoring

**Manage the host via RESTAPI**



# Configuration Management

## High availability manager

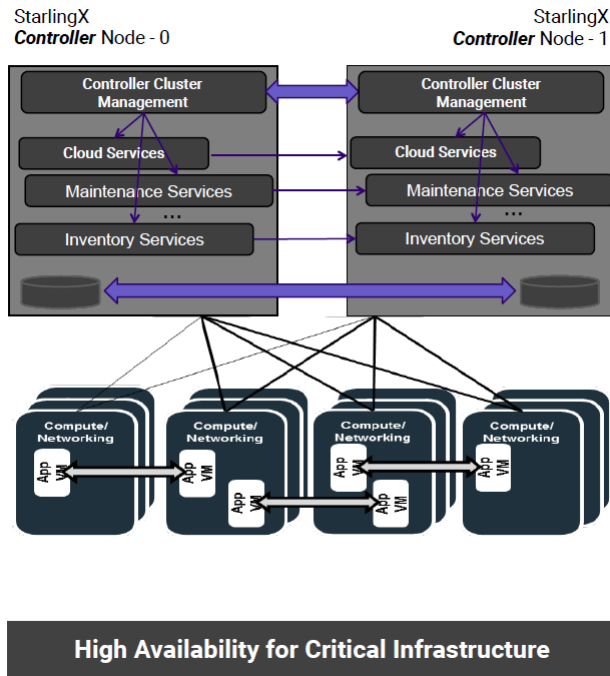
- Redundancy model can be N+M or N across multiple nodes
- Currently 1+1 HA Controller Cluster

## Uses multiple messaging paths to avoid split-brain communication failures

- Up to 3 independent communication paths
- LAG can also be configured for multi-link protection of each path
- Messages are authenticated using HMAC
- SHA-512 if configured / enabled on an interface- by-interface basis

## Active or passive monitoring of services

## Allows for specifying the impact of a service failure



# Software Management

**Automated deploy of software updates for security and/or new functionality**

## **Integrated end-to-end rolling upgrade solution**

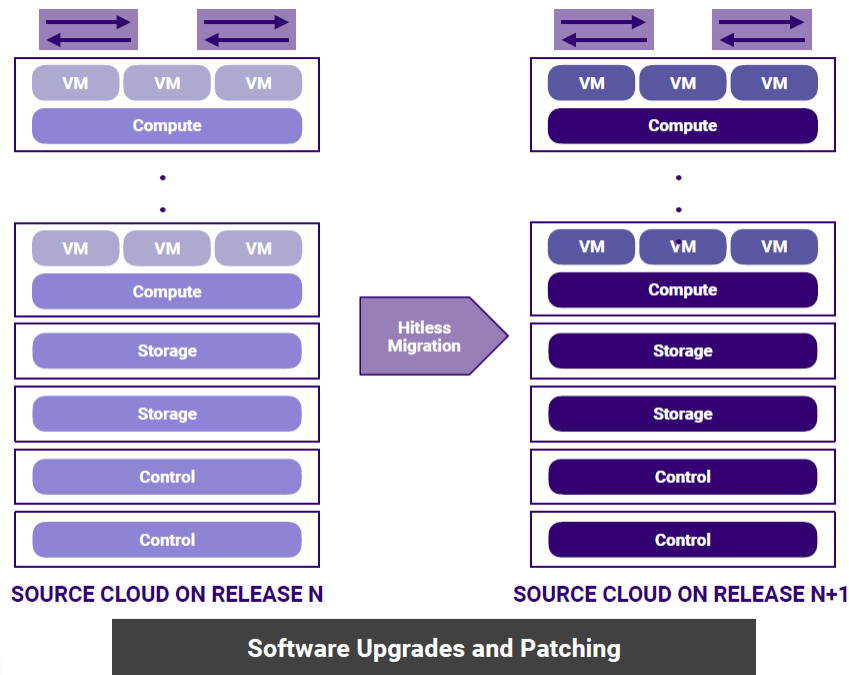
- Automated, low number of steps
- No additional hardware required for upgrade
- Rolling upgrade across nodes

## **In-service and reboot required patches supported**

- Reboot required for kernel replacement etc.
- VM live migration is used for patches that require reboot

## **Manages upgrades of all software**

- Host OS changes
- New / upgraded StarlingX service software
- New / upgraded OpenStack software





# Fault Management

## Framework for infrastructure services via API

- Set, clear and query customer alarms
- Generate customer logs for significant events

## Maintains an Active Alarm List

## Provides REST API to query alarms and events


## Support for alarm suppression

## Operator alarms

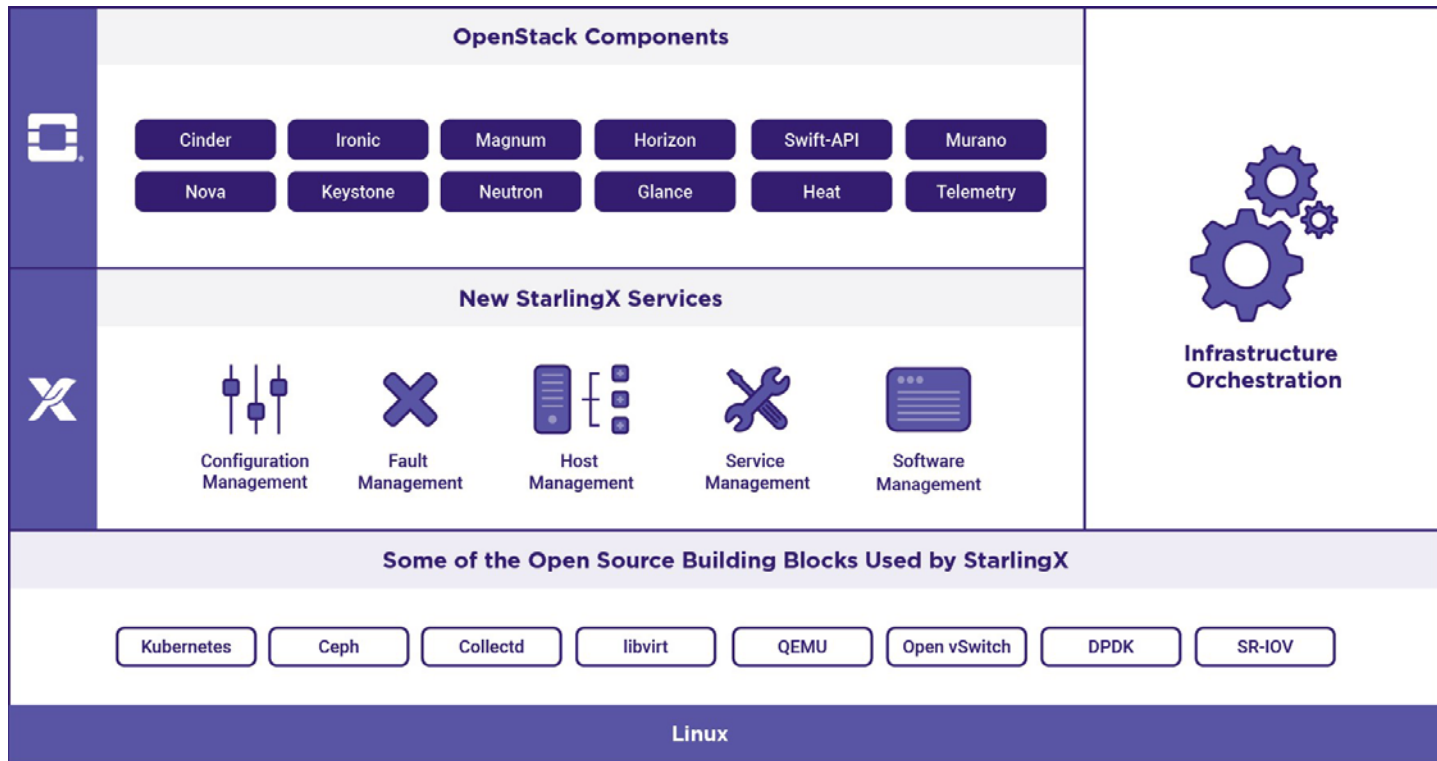
- On platform nodes and resources
- On hosted virtual resources

## Operator logs - Event List

- Logging of set/Clears of alarms
- Related to platform nodes and resources
- Related to hosted virtual resources



# Current Architecture



# Next Generation Container Architecture

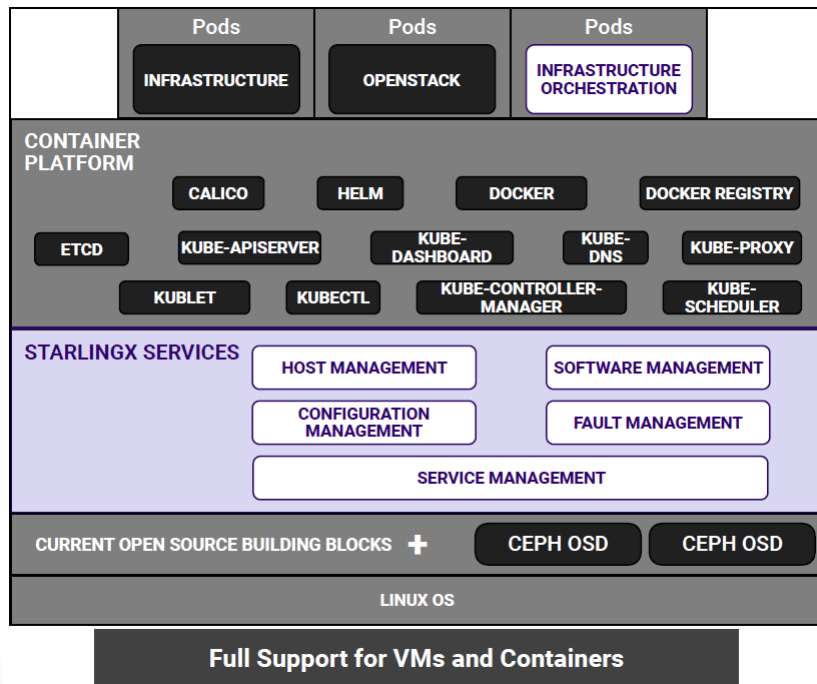
## StarlingX is evolving to

- Running OpenStack containerized
- On top of a bare metal Kubernetes cluster
- With OpenStack Helm managing the lifecycle of the OpenStack cluster

## With Kubernetes Cluster initially supporting

- Docker runtime
- Calico CNI plugin
- CEPH as persistent storage backend
- HELM as the package manager
- Local Docker Image Registry

**Along with Kubernetes cluster available for non-OpenStack end user applications**



# The Road to the Edge

## Build it yourself from open source components

- Building blocks need refinement
- Time consuming
- Gaps to fill

## Use StarlingX

- New services provide improved manageability for the platform and high availability for your applications to meet Edge Cloud requirements
- Tested and available as a complete stack
- Mission-ready for your applications

# StarlingX Community & Activities

- StarlingX had the first release on October 24. The release includes 1,329 commits from 84 contributors, including developers representing Wind River, Intel, 99Cloud, China Unionpay, Suse, Fujitsu, and NEC, among others.
- After the first release the community is focusing on new functionality as well as making their tools and processes more efficient. As part of this work StarlingX is working on containerization it's infrastructure and improving the build and deployment experience both for users and developers.
- The first StarlingX Technical Steering Committee (TSC) has been formed with 8 members from 6 companies: Ana Cunha (Ericsson), Brent Rowsell (Wind River), Curtis Collicutt (INTERdynamix), Dean Troyer (Intel), Ian Jolliffe (Wind River), Miguel Lavallo (Huawei), Saul Wold (Intel), Shuquan Huang (99cloud)

# Sub-project Structure

## Main sub-projects

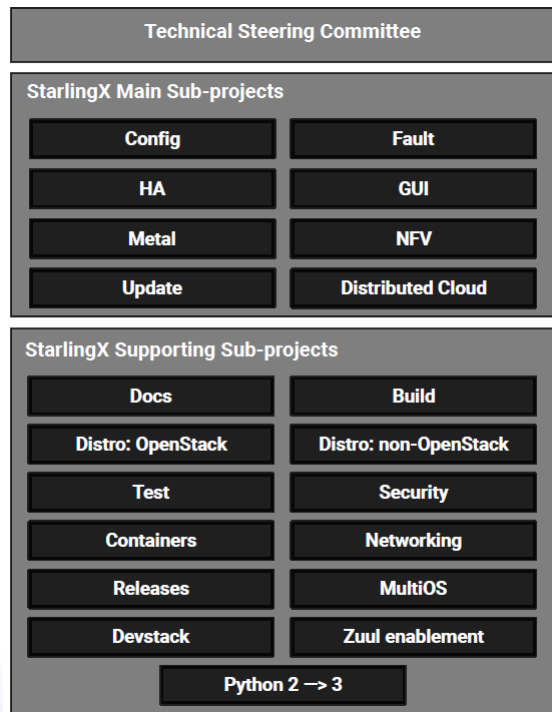
- New functionality and services

## Supporting sub-projects

- Supporting services, test and infrastructure

## Sub-project team structure

- 1 Team Lead
- 1 Project Lead
- Core Reviewers
- Contributors





# Get Involved

**Code and documentation are available through git**

- [git.starlingx.io](https://git.starlingx.io)

**Apache 2 license**

**IRC: #starlingx@Freenode**

**Mailing List for daily discussions**

- <http://lists.starlingx.io/cgi-bin/mailman/listinfo/starlingx-discuss>

# Where to Contribute?

## Bugs are tracked in Launchpad

- <https://bugs.launchpad.net/starlingx>

## New ideas are introduced in the specs repository

- <https://git.openstack.org/cgit/openstack/stx-specs/>

## Design and implementation work is tracked in StoryBoard

- [https://storyboard.openstack.org/#!/project\\_group/86](https://storyboard.openstack.org/#!/project_group/86)

# Communication

**#starlingx@Freenode, IRC channel for online discussions**

**Mailing Lists: [lists.starlingx.io](https://lists.starlingx.io)**

**Email: [info@starlingx.io](mailto:info@starlingx.io)**

**Weekly meetings:**

- Zoom calls
- <https://wiki.openstack.org/wiki/Starlingx/Meetings>

**Twitter handle: @StarlingX**

The background is a deep blue with abstract digital elements. In the top left, there are interlocking gears. A network of glowing blue nodes connected by lines is visible in the upper right. On the right side, there is a profile of a human head composed of a grid of blue dots. In the bottom left, there are wavy, glowing lines in shades of orange and purple, overlaid with a network of white lines.

# THANKS!